



FAT

The **File Allocation Table (FAT)** file system was introduced with DOS v1.0 (and possibly CP/M). Supposedly written by Bill Gates, FAT is a very simple file system -- nothing more than a singly-linked list of clusters in a gigantic table. A FAT file system uses very little memory (unless the OS caches the whole allocation table in memory) and is one of, if not the, most basic file system in use today.

Contents [hide]

- 1 Overview
 - 1.1 FAT 12
 - 1.2 FAT 16
 - 1.3 FAT 32
 - 1.4 ExFAT
 - 1.5 VFAT
- 2 Implementation Details
 - 2.1 Boot Record
 - 2.1.1 BPB (BIOS Parameter Block)
 - 2.1.2 Extended Boot Record
 - 2.1.2.1 FAT 12 and FAT 16
 - 2.1.2.2 FAT 32
 - 2.2 FSInfo Structure (FAT32 only)
 - 2.3 File Allocation Table
 - 2.3.1 FAT 12
 - 2.3.2 FAT 16
 - 2.3.3 FAT 32
 - 2.4 Directories
 - 2.4.1 Standard 8.3 format
 - 2.4.2 Long File Names
- 3 Programming Guide
 - 3.1 Reading the Boot Sector
 - 3.2 Reading Directories
 - 3.3 Following Cluster Chains
- 4 Creating a fresh FAT filesystem
- 5 See Also
 - 5.1 Threads
 - 5.2 External Links

Overview

There are several different versions of the FAT file system. Each version was designed for a different size of storage media.

FAT 12

FAT 12 was designed for floppy disks and can manage a maximum size of 16 megabytes because it uses 12 bits to address the clusters.

FAT 16

FAT 16 was designed for early hard disks and could handle a maximum size of 64K clusters * the cluster size. The larger the hard disk, the larger the cluster size would be, which leads to large amounts of "slack space" on the disk.

FAT 32

FAT 32 was introduced to us by Windows95-B and Windows98. FAT32 solved some of FAT's problems. No more 64K max clusters! FAT32 is slightly misnamed, as the top 4 bits of the 32 bit cluster number are reserved, and were never used. If you want to call it FAT28 instead, then as the name suggests, the filesystem can handle a maximum of 256M clusters per partition. This enables very large hard disks to still maintain reasonably small cluster sizes and thus reduce slack space between files.

ExFAT

Main article: [ExFAT](#)

ExFAT is the filesystem used on SDXC cards, created by Microsoft. It is basically FAT32 with actually 32 bits per FAT entry, with minor extensions.

Microsoft has published the official specification at <https://docs.microsoft.com/en-us/windows/win32/fileio/exfat-specification> .

VFAT

VFAT is an extension to the FAT file system that has the ability to use long filenames (up to 255 characters). First introduced by Windows 95, it uses a "kludge" whereby long filenames are marked with a "volume label" attribute and filenames are subsequently stored in 11 byte chunks in sequential directory entries. (This is a bit of an oversimplification, but close enough).

Implementation Details

The FAT file system views the storage media as a flat array of clusters. If the physical media does not address its data as a flat list of sectors (really old hard disks and floppy disks) then the cluster numbers will need to be translated before being sent to the disk. The storage media is organized into three basic areas.

- The boot record
- The File Allocation Table (FAT)
- The directory and data area

Boot Record

The boot record occupies one sector, and is always placed in logical sector number zero of the "partition". If the media is not divided into partitions, then this is the beginning of the media. This is the easiest sector on the partition for the computer to locate when it is loaded. If the storage media is partitioned (such as a hard disk), then the beginning of the actual media contains an [MBR \(x86\)](#) or other form of partition information. In this case each partition's first sector holds a [Volume Boot Record](#).

BPB (BIOS Parameter Block)

Filesystems

Virtual Filesystems

[VFS](#)

Disk Filesystems

FAT [12/16/32](#), [VFAT](#)

[Ext 2/3/4](#)

[LEAN](#)

[HPFS](#)

[NTFS](#)

[HFS](#)

[HFS+](#)

[MFS](#)

[ReiserFS](#)

[FFS \(Amiga\)](#)

[FFS \(BSD\)/UFS](#)

[BeFS](#)

[BFS](#)

[XFS](#)

[SFS](#)

[ZFS](#)

[USTAR](#)

CD/DVD Filesystems

[ISO 9660](#)

[Joliet](#)

[UDF](#)

Network Filesystems

[NFS](#)

[RFS](#)

[AFS](#)

Flash Filesystems

[JFFS2](#)

[YAFFS](#)

The boot record contains both code and data, mixed together. The data that isn't code is known as the BPB.

Offset (decimal)	Offset (hex)	Size (in bytes)	Meaning
0	0x00	3	The first three bytes EB 3C 90 disassemble to JMP SHORT 3C NOP. (The 3C value may be different.) The reason for this is to jump over the disk format information (the BPB and EBPB). Since the first sector of the disk is loaded into ram at location 0x0000:0x7c00 and executed, without this jump, the processor would attempt to execute data that isn't code. Even for non-bootable volumes, code matching this pattern (or using the E9 jump opcode) is required to be present by both Windows and OS X. To fulfil this requirement, an infinite loop can be placed here with the bytes EB FE 90.
3	0x03	8	OEM identifier. The first 8 Bytes (3 - 10) is the version of DOS being used. The next eight Bytes 29 3A 63 7E 2D 49 48 and 43 read out the name of the version. The official FAT Specification from Microsoft says that this field is really meaningless and is ignored by MS FAT Drivers, however it does recommend the value "MSWIN4.1" as some 3rd party drivers supposedly check it and expect it to have that value. Older versions of dos also report MSDOS5.1, linux-formatted floppy will likely to carry "mkdofs" here, and FreeDOS formatted disks have been observed to have "FRDOS5.1" here. If the string is less than 8 bytes, it is padded with spaces.
11	0x0B	2	The number of Bytes per sector (remember, all numbers are in the little-endian format).
13	0x0D	1	Number of sectors per cluster.
14	0x0E	2	Number of reserved sectors. The boot record sectors are included in this value.
16	0x10	1	Number of File Allocation Tables (FAT's) on the storage media. Often this value is 2.
17	0x11	2	Number of directory entries (must be set so that the root directory occupies entire sectors).
19	0x13	2	The total sectors in the logical volume. If this value is 0, it means there are more than 65535 sectors in the volume, and the actual count is stored in the Large Sector Count entry at 0x20.
21	0x15	1	This Byte indicates the media descriptor type 🗑️
22	0x16	2	Number of sectors per FAT. FAT12/FAT16 only.
24	0x18	2	Number of sectors per track.
26	0x1A	2	Number of heads or sides on the storage media.
28	0x1C	4	Number of hidden sectors. (i.e. the LBA of the beginning of the partition.)
32	0x20	4	Large sector count. This field is set if there are more than 65535 sectors in the volume, resulting in a value which does not fit in the <i>Number of Sectors</i> entry at 0x13.

Note: the "geometry" of the media (sectors per track, heads, and perhaps the number of bytes in a sector) is not necessarily known correctly by the program that originally formats the media. Also, if the media is moved (from the computer that formatted it) to another machine with a different BIOS -- then the new BIOS may specify a different geometry for the same media. So it is generally a very bad idea to trust the "SPT" or "heads" numbers. Get them from the BIOS instead, if possible.

Note2: many of the values in the BPB are not correctly "aligned". That is, word-sized values are not stored on word ("even" address) boundaries. On some architectures, accessing misaligned words may cause the code to crash. Making a copy of the BPB (somewhere else in memory and shifted up one byte) may solve the problem.

Extended Boot Record

The extended boot record information comes right after the BPB. The data at the beginning is known as the EBPB. It contains different information depending on whether this partition is a FAT 12, FAT 16, or FAT 32 filesystem. Immediately following the EBPB is the actual boot code, then the standard 0xAA55 boot signature, to fill out the 512-byte boot sector. Offsets shows are from the start of the standard boot record.

FAT 12 and FAT 16

Offset (decimal)	Offset (hexadecimal)	Length (in bytes)	Meaning
36	0x024	1	Drive number. The value here should be identical to the value returned by BIOS interrupt 0x13, or passed in the DL register; i.e. 0x00 for a floppy disk and 0x80 for hard disks. This number is useless because the media is likely to be moved to another machine and inserted in a drive with a different drive number.
37	0x025	1	Flags in Windows NT. Reserved otherwise.
38	0x026	1	Signature (must be 0x28 or 0x29).
39	0x027	4	VolumeID 'Serial' number. Used for tracking volumes between computers. You can ignore this if you want.
43	0x02B	11	Volume label string. This field is padded with spaces.
54	0x036	8	System identifier string. This field is a string representation of the FAT file system type. It is padded with spaces. The spec says never to trust the contents of this string for any use.
62	0x03E	448	Boot code.
510	0x1FE	2	Bootable partition signature 0xAA55.

FAT 32

Offset (decimal)	Offset (hexadecimal)	Length (in bytes)	Meaning
36	0x024	4	Sectors per FAT. The size of the FAT in sectors.
40	0x028	2	Flags.
42	0x02A	2	FAT version number. The high byte is the major version and the low byte is the minor version. FAT drivers should respect this field.
44	0x02C	4	The cluster number of the root directory. Often this field is set to 2.
48	0x030	2	The sector number of the FSInfo structure.
50	0x032	2	The sector number of the backup boot sector.
52	0x034	12	Reserved. When the volume is formatted these bytes should be zero.

64	0x040	1	Drive number. The values here are identical to the values returned by the BIOS interrupt 0x13. 0x00 for a floppy disk and 0x80 for hard disks.
65	0x041	1	Flags in Windows NT. Reserved otherwise.
66	0x042	1	Signature (must be 0x28 or 0x29).
67	0x043	4	Volume ID 'Serial' number. Used for tracking volumes between computers. You can ignore this if you want.
71	0x047	11	Volume label string. This field is padded with spaces.
82	0x052	8	System identifier string. Always "FAT32 ". The spec says never to trust the contents of this string for any use.
90	0x05A	420	Boot code.
510	0x1FE	2	Bootable partition signature 0xAA55.

FSInfo Structure (FAT32 only)

Offset (decimal)	Offset (hexadecimal)	Length (in bytes)	Meaning
0	0x0	4	Lead signature (must be 0x41615252 to indicate a valid FSInfo structure)
4	0x4	480	Reserved, these bytes should never be used
484	0x1E4	4	Another signature (must be 0x61417272)
488	0x1E8	4	Contains the last known free cluster count on the volume. If the value is 0xFFFFFFFF, then the free count is unknown and must be computed. However, this value might be incorrect and should at least be range checked (<= volume cluster count)
492	0x1EC	4	Indicates the cluster number at which the filesystem driver should start looking for available clusters. If the value is 0xFFFFFFFF, then there is no hint and the driver should start searching at 2. Typically this value is set to the last allocated cluster number. As the previous field, this value should be range checked.
496	0x1F0	12	Reserved
508	0x1FC	4	Trail signature (0xAA550000)

File Allocation Table

The File Allocation Table (FAT) is a table stored on the storage media that indicates the status and location of all data clusters that are on the disk. It can be considered the "table of contents" of a disk. The cluster may be available for use, it may be reserved by the operating system, it may be unavailable due to a bad sector on the disk, or it may be in use by a file. The clusters of a file need not be right next to each other on the disk. In fact it is likely that they are scattered widely throughout the disk. The FAT allows the operating system to follow the "chain" of clusters in a file.

FAT 12

FAT 12 uses 12 bits to address the clusters on the disk. Each 12 bit entry in the FAT points to the next cluster of a file on the disk. Given a valid cluster number, here is how you extract the value of the next cluster in the cluster chain:

```
unsigned char FAT_table[sector_size];
unsigned int fat_offset = active_cluster + (active_cluster / 2); // multiply by 1.5
unsigned int fat_sector = first_fat_sector + (fat_offset / sector_size);
unsigned int ent_offset = fat_offset % sector_size;

//at this point you need to read from sector "fat_sector" on the disk into "FAT_table".

unsigned short table_value = *(unsigned short*)&FAT_table[ent_offset];

if(active_cluster & 0x0001)
    table_value = table_value >> 4;
else
    table_value = table_value & 0x0FFF;

//the variable "table_value" now has the information you need about the next cluster in the chain.
```

If "table_value" is greater than or equal to (>=) 0xFF8 then there are no more clusters in the chain. This means that the whole file has been read. If "table_value" equals (==) 0xFF7 then this cluster has been marked as "bad". "Bad" clusters are prone to errors and should be avoided. If "table_value" is not one of the above cases then it is the cluster number of the next cluster in the file.

Since FAT12 uses an entry size that is not evenly divisible by 8 bits, figuring out how to interpret the FAT on a little-endian machine can be slightly confusing. Consider two entries of 0x123 and 0x456 back-to-back. On a little-endian machine, the first byte of the first entry is the bottom two nibbles (0x23) and the highest nibble goes into the bottom nibble of the second byte (0x71). Since the next entry is now starting mid-byte, only the lowest nibble can fit in the byte (0x67) and the two highest nibbles go into the next byte (0x45). Therefore the 2 entries back-to-back look like this: 0x23 0x61 0x45.

FAT 16

FAT 16 uses 16 bits to address the clusters on the disk. Because of this, it is much easier to extract the values out of a 16 bit File Allocation Table. Here is how it is done:

```
unsigned char FAT_table[sector_size];
unsigned int fat_offset = active_cluster * 2;
unsigned int fat_sector = first_fat_sector + (fat_offset / sector_size);
unsigned int ent_offset = fat_offset % sector_size;

//at this point you need to read from sector "fat_sector" on the disk into "FAT_table".

unsigned short table_value = *(unsigned short*)&FAT_table[ent_offset];

//the variable "table_value" now has the information you need about the next cluster in the chain.
```

If "table_value" is greater than or equal to (>=) 0xFFF8 then there are no more clusters in the chain. This means that the whole file has been read. If "table_value" equals (==) 0xFFF7 then this cluster has been marked as "bad". "Bad" clusters are prone to errors and should be avoided. If "table_value" is not one of the above cases then it is the cluster

number of the next cluster in the file.

FAT 32

FAT 32 uses 28 bits to address the clusters on the disk. Yes, that is right. FAT 32 only uses 28 of its 32 bits. The highest 4 bits are reserved. This means that they should be ignored when read and unchanged when written. Besides this small detail, extracting a value from a 32 bit FAT is almost identical to the same operation on a 16 bit FAT:

```
unsigned char FAT_table[sector_size];
unsigned int fat_offset = active_cluster * 4;
unsigned int fat_sector = first_fat_sector + (fat_offset / sector_size);
unsigned int ent_offset = fat_offset % sector_size;

//at this point you need to read from sector "fat_sector" on the disk into "FAT_table".

//remember to ignore the high 4 bits.
unsigned int table_value = *(unsigned int*)&FAT_table[ent_offset] & 0x0FFFFFFF;

//the variable "table_value" now has the information you need about the next cluster in the chain.
```

If "table_value" is greater than or equal to (\geq) 0x0FFFFFF8 then there are no more clusters in the chain. This means that the whole file has been read. If "table_value" equals ($=$) 0x0FFFFFF7 then this cluster has been marked as "bad". "Bad" clusters are prone to errors and should be avoided. If "table_value" is not one of the above cases then it is the cluster number of the next cluster in the file.

Directories

A directory entry simply stores the information needed to know where a file's data or a folder's children are stored on the disk. It also holds information such as the entry's name, size, and creation time. There are two types of directories in a FAT file system. Standard 8.3 directory entries, which appear on all FAT file systems, and Long File Name directory entries which are optionally present to allow for longer file names.

Standard 8.3 format

Offset (in bytes)	Length (in bytes)	Meaning						
0	11	8.3 file name. The first 8 characters are the name and the last 3 are the extension.						
11	1	Attributes of the file. The possible attributes are: <div><div>READ_ONLY=0x01 HIDDEN=0x02 SYSTEM=0x04 VOLUME_ID=0x08 DIRECTORY=0x10 ARCHIVE=0x20 LFN=READ_ONLY HIDDEN SYSTEM VOLUME_ID</div><div>(LFN means that this entry is a long file name entry)</div></div>						
12	1	Reserved for use by Windows NT.						
13	1	Creation time in tenths of a second. Range 0-199 inclusive. Based on simple tests, Ubuntu16.10 stores either 0 or 100 while Windows7 stores 0-199 in this field.						
14	2	The time that the file was created. Multiply Seconds by 2. <table><tr><td>Hour</td><td>5 bits</td></tr><tr><td>Minutes</td><td>6 bits</td></tr><tr><td>Seconds</td><td>5 bits</td></tr></table>	Hour	5 bits	Minutes	6 bits	Seconds	5 bits
Hour	5 bits							
Minutes	6 bits							
Seconds	5 bits							
16	2	The date on which the file was created. <table><tr><td>Year</td><td>7 bits</td></tr><tr><td>Month</td><td>4 bits</td></tr><tr><td>Day</td><td>5 bits</td></tr></table>	Year	7 bits	Month	4 bits	Day	5 bits
Year	7 bits							
Month	4 bits							
Day	5 bits							
18	2	Last accessed date. Same format as the creation date.						
20	2	The high 16 bits of this entry's first cluster number. For FAT 12 and FAT 16 this is always zero.						
22	2	Last modification time. Same format as the creation time.						
24	2	Last modification date. Same format as the creation date.						
26	2	The low 16 bits of this entry's first cluster number. Use this number to find the first cluster for this entry.						
28	4	The size of the file in bytes.						

Long File Names

Long file name entries *always* have a regular 8.3 entry to which they belong. The long file name entries are always placed immediately before their 8.3 entry. Here is the format of a long file name entry.

Offset (in bytes)	Length (in bytes)	Meaning
0	1	The order of this entry in the sequence of long file name entries. This value helps you to know where in the file's name the characters from this entry should be placed.
1	10	The first 5, 2-byte characters of this entry.
11	1	Attribute. Always equals 0x0F. (the long file name attribute)
12	1	Long entry type. Zero for name entries.
13	1	Checksum generated of the short file name when the file was created. The short filename can change without changing the long filename in cases where the partition is mounted on a system which does not support long filenames.
14	12	The next 6, 2-byte characters of this entry.

26	2	Always zero.
28	4	The final 2, 2-byte characters of this entry.

Here is an example of what a regular 8.3 entry with one long file name entry preceding it might look like in a hex editor:

```
41 62 00 69 00 6E 00 00 00 FF FF 0F 00 7F FF FF FF FF FF FF FF FF FF FF FF 00 00 FF FF FF FF
42 49 4E 20 20 20 20 20 20 20 20 10 00 00 F7 01 D5 38 D5 38 00 00 F7 01 D5 38 03 00 00 00 00 00
```

And in a text editor:

```
Ab.i.n.....
BIN      .....8.8.....
```

The first line is the long file name entry (the second line is the regular 8.3 entry). The very first byte (41) tells us two important pieces of information. First, the one (01) tells us that this is the first long file name entry for the regular 8.3 entry. Second the forty (40) part tells us that this is also the *last* long file name entry for this regular 8.3 entry. The next 10 bytes spell out the first part of the long file name. In this case they read:

```
b 00 i 00 n 00 00 00 FF FF
```

Notice that each character is two bytes long and that the name is null terminated. The two FF's at the end are the padding at the end of the long file name. This is also what the other FF's in the long file name entry are. The final important thing to notice about the long file name entry is it's attribute byte at offset 11. the 0x0F attribute allows us to verify that this is indeed a long file name entry.

Programming Guide

This section is intended to give you information about common functions that are preformed on a FAT file system. Edit: <plug>If you are focusing on FAT32 only, there is an article for that: [User:requimrar/FAT32](#) that you might want to check out.</plug>

Reading the Boot Sector

The Boot Sector is always placed at logical sector number zero. You can either read the boot sector into an array and access it's members that way or you can read it into a structure and access it through the structure. Either way, the values in the boot sector need to be readily available in order to do much of anything with a FAT file system.

Here is an example of some boot sector structures in C.

```
typedef struct fat_extBS_32
{
    //extended fat32 stuff
    unsigned int    table_size_32;
    unsigned short  extended_flags;
    unsigned short  fat_version;
    unsigned int    root_cluster;
    unsigned short  fat_info;
    unsigned short  backup_BS_sector;
    unsigned char   reserved_0[12];
    unsigned char   drive_number;
    unsigned char   reserved_1;
    unsigned char   boot_signature;
    unsigned int    volume_id;
    unsigned char   volume_label[11];
    unsigned char   fat_type_label[8];
}__attribute__((packed)) fat_extBS_32_t;

typedef struct fat_extBS_16
{
    //extended fat12 and fat16 stuff
    unsigned char   bios_drive_num;
    unsigned char   reserved1;
    unsigned char   boot_signature;
    unsigned int    volume_id;
    unsigned char   volume_label[11];
    unsigned char   fat_type_label[8];
}__attribute__((packed)) fat_extBS_16_t;

typedef struct fat_BS
{
    unsigned char   bootjmp[3];
    unsigned char   oem_name[8];
    unsigned short  bytes_per_sector;
    unsigned char   sectors_per_cluster;
    unsigned short  reserved_sector_count;
    unsigned char   table_count;
    unsigned short  root_entry_count;
    unsigned short  total_sectors_16;
    unsigned char   media_type;
    unsigned short  table_size_16;
    unsigned short  sectors_per_track;
    unsigned short  head_side_count;
    unsigned int    hidden_sector_count;
    unsigned int    total_sectors_32;

    //this will be cast to it's specific type once the driver actually knows what type of FAT this is.
    unsigned char   extended_section[54];
}__attribute__((packed)) fat_BS_t;
```

Important pieces of information that can be extracted from the boot sector include:

Total sectors in volume (including VBR):

```
total_sectors = (fat_boot->total_sectors_16 == 0)? fat_boot->total_sectors_32 : fat_boot->total_sectors_16;
```

FAT size in sectors:

```
fat_size = (fat_boot->table_size_16 == 0)? fat_boot_ext_32->table_size_16 : fat_boot->table_size_16;
```

The size of the root directory (unless you have FAT32, in which case the size will be 0):

```
root_dir_sectors = ((fat_boot->root_entry_count * 32) + (fat_boot->bytes_per_sector - 1)) / fat_boot->bytes_per_sector;
```

This calculation will round up. 32 is the size of a FAT directory in bytes.

The first data sector (that is, the first sector in which directories and files may be stored):

```
first_data_sector = fat_boot->reserved_sector_count + (fat_boot->table_count * fat_size) + root_dir_sectors;
```

The first sector in the File Allocation Table:

```
first_fat_sector = fat_boot->reserved_sector_count;
```

The total number of data sectors:

```
data_sectors = fat_boot->total_sectors - (fat_boot->reserved_sector_count + (fat_boot->table_count * fat_size) + root_dir_sectors)
```

The total number of clusters:

```
total_clusters = data_sectors / fat_boot->sectors_per_cluster;
```

This rounds down.

The FAT type of this file system:

```
if(total_clusters < 4085)
{
    fat_type = FAT12;
}
else if(total_clusters < 65525)
{
    fat_type = FAT16;
}
else if (total_clusters < 268435445)
{
    fat_type = FAT32;
}
else
{
    fat_type = ExFAT;
}
```

Reading Directories

The first step in reading directories is finding and reading the root directory. On a FAT 12 or FAT 16 volumes the root directory is at a fixed position immediately after the File Allocation Tables:

```
first_root_dir_sector = first_data_sector - root_dir_sectors;
```

In FAT32, root directory appears in data area on given cluster and can be a cluster chain.

```
root_cluster_32 = extBS_32->root_cluster;
```

For each given cluster number we can calculate the first sector of it (relative to the partition's offset):

```
first_sector_of_cluster = ((cluster - 2) * fat_boot->sectors_per_cluster) + first_data_sector;
```

After the correct cluster has been loaded into memory, the next step is to read and parse all of the entries in it. Each entry is 32 bytes long. For each 32 byte entry this is the flow of execution:

1. If the first byte of the entry is equal to 0 then there are no more files/directories in this directory. FirstByte==0, finish. FirstByte!=0, goto 2.
2. If the first byte of the entry is equal to 0xE5 then the entry is unused. FirstByte==0xE5, goto 8, FirstByte!=0xE5, goto 3.
3. Is this entry a long file name entry? If the 11'th byte of the entry equals 0x0F, then it is a long file name entry. Otherwise, it is not. 11thByte==0x0F, goto 4. 11thByte!=0x0F, goto 5.
4. Read the portion of the long filename into a temporary buffer. Goto 8.
5. Parse the data for this entry using the table from further up on this page. It would be a good idea to save the data for later. Possibly in a virtual file system structure. goto 6
6. Is there a long file name in the temporary buffer? Yes, goto 7. No, goto 8
7. Apply the long file name to the entry that you just read and clear the temporary buffer. goto 8
8. Increment pointers and/or counters and check the next entry. (goto number 1)

This process should be repeated until all of the entries have been read from the cluster. You should then check to see if there is another cluster following this one in the cluster chain or if this is the last cluster in the chain. See [section below](#) and [FAT](#) section for more information. You should do the above process for each cluster in the chain, following it

until there are no more clusters left in the chain. Then you can check if any of the entries that you just read are directories. If they are they should each be read in the same way starting with their first cluster number which is stored in the entry.

Following Cluster Chains

There are two basic steps to following cluster chains. The first step is to find out if there is another "link" (cluster) following the current one in the chain. The second step is to actually use the value read from the FAT to read the next sector. Here is the basic idea:

1. Extract the value from the FAT for the `_current_` cluster. (Use the previous section on the File Allocation Table for details on how exactly to extract the value.) goto number 2
2. Is this cluster marked as the last cluster in the chain? (again, see the above section for more details) Yes, goto number 4. No, goto number 3
3. Read the cluster represented by the extracted value and return for more directory parsing.
4. The end of the cluster chain has been found. Our work here is finished. :)

Creating a fresh FAT filesystem

Typically during development you want to create a disk image with a FAT filesystem. There are two common approaches for this, either by using a utility that works directly on images, or by using a [Loopback Device](#) and using the OS' own driver to work on the image. A less common alternative is to have an actual disk in your drive.

The most buildscript-friendly tool is [MTools](#) - which can do all operations directly on a disk image using the `-i` argument and supplies every DOS command related to files in this fashion, only prefixed with an `m.` It can also use a configuration file to access drives in their DOS fashion, allowing you to use for instance `A:` and `C:` as actual drives. The tool can be built out of the box for Windows and is included in many a linux package manager.

Linux-only developers can, often with a bit of `sudo` and permission magic, automate the [Loopback Device](#) in combination with `mkdosfs` or `mkfs.vfat` as well as partition editing. This method is less portable as the commands often can't be reused outside of Linux. Several developers also make the error of passing `-F` to `mkdosfs` in an attempt to choose a FAT size, which often has the effect of creating a corrupt filesystem since the result doesn't follow [the official rule for FAT sizes anymore](#).

Windows users can make use of [VFD](#) for loopback devices. It comes with a GUI, but at the cost of not being properly automatable in a script.

See Also

Threads

- from raw bits to directory listing (code posted) in the [forum](#)
- Public Domain FAT32 code in the [forum](#)
- FAT12/FAT16 bootsector code in the [forum](#)

External Links

- [FAT32 File System Specification](#) - from Cottontail OS Development Library
- [FAT32 File System Specification](#) - from Microsoft (documentation, but in tutorial style with many code examples)
- [About an error in the above specification](#)
- <http://scottie.20m.com/fat.htm>
- http://www.maverick-os.dk/FileSystemFormats/FAT12_FileSystem.html
- <http://www.pjrc.com/tech/8051/ide/fat32.html>
- [Intro into Sectors and Addressing](#)
- http://elm-chan.org/fsw/ff/00index_e.html - simple (V)FAT12/16/32 read/write library with good documentation
- <http://gitorious.org/unix-stuff/fat-util> - Utility to read, remove and extract files on FAT12, 16 and 32
- <http://www.larwe.com/zws/products/dosfs/index.html> - fat12/16/32 compatible fs driver
- <http://www.isdaman.com/alsos/protocols/fats/nowhere/FAT.HTM>

Category: [Filesystems](#)

This page was last modified on 9 October 2019, at 12:56.

This page has been accessed 239,183 times.

[Privacy policy](#) [About OSDev Wiki](#) [Disclaimers](#)

